

1. There is a possibility of using S3 (Simple Storage Service). The most popular one is Amazon S3, but we'd prefer domestic alternatives. We have examples from companies like Mail.ru and others. Essentially, they are structured like a file system, except for a few differences: higher security due to using copies, access via the HTTP protocol, and scalability. A traditional database won't work because it's used for structured data, which is not our case. Since all S3 options are paid and not cheap, it's likely we need to use our own server for data storage. The open question is how to organize it similar to S3 because we need the ability to work over HTTP. Perhaps there's a way to bypass the need for HTTP, which I'm not aware of; this is an open question.
2. Fast access is probably impossible. Snapshots can be gigabytes in size, which takes several minutes to load. Archiving -> transfer -> unarchiving is a minus right away because it's more complicated, and there's no benefit since archiving and unarchiving also take time, often not insignificant. So in any case, the first opening of a snapshot is like going for coffee for the user. After that, we can keep it in a cache. But that's also a hassle, and there's not much benefit. The cache can grow significantly. Another problem is that I couldn't open a .tif file in the browser, neither with standard methods nor with third-party ones. Of course, you can compress everything to jpg beforehand, but then it's unclear why these precise snapshots are needed at all.
3. Integration is possible through caching so that other services can open previously viewed snapshots faster. The problems are the same as mentioned above.

For now, I see the solution as follows: organize the upload and download using our server as storage. Users can open downloaded snapshots in QGIS or even in the built-in editor.

More details are needed about the purpose of these snapshots for the user. It's somewhat clear on the surface, but I'd like more details.